

Introduction to Probabilities and Statistics

Arnaud Legrand and Jean-Marc Vincent

Scientific Methodology and Performance Evaluation
M2R MOSIG, Grenoble, September-December 2015

Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Probabilities

- Using probabilities enables to model **uncertainty** that may result of **incomplete information** or **imprecise measurements**

Probabilities

- Using probabilities enables to model **uncertainty** that may result of **incomplete information** or **imprecise measurements**

A **random variable** (or stochastic variable) is, roughly speaking, a variable whose value results from a measurement (or an observation)

You can think of it as a **small box**:

- Every time you open the box, you get a different value.
- I will use this box analogy throughout the whole lecture and I encourage you to ask yourself what the box can be in your own studies

Probabilities

- Using probabilities enables to model **uncertainty** that may result of **incomplete information** or **imprecise measurements**

A **random variable** (or stochastic variable) is, roughly speaking, a variable whose value results from a measurement (or an observation)

You can think of it as a **small box**:

- Every time you open the box, you get a different value.
- I will use this box analogy throughout the whole lecture and I encourage you to ask yourself what the box can be in your own studies
- Formally a **probability space** is defined by (Ω, \mathcal{F}, P) where:
 - Ω , the **sample space**, is the set of all possible **outcomes**
 - E.g., all the possible combinations of your DNA with the one of your {girl|boy}friend
 - You may or may not be able to observe directly the outcome.

Probabilities

- Using probabilities enables to model **uncertainty** that may result of **incomplete information** or **imprecise measurements**

A **random variable** (or stochastic variable) is, roughly speaking, a variable whose value results from a measurement (or an observation)

You can think of it as a **small box**:

- Every time you open the box, you get a different value.
- I will use this box analogy throughout the whole lecture and I encourage you to ask yourself what the box can be in your own studies
- Formally a **probability space** is defined by (Ω, \mathcal{F}, P) where:
 - Ω , the **sample space**, is the set of all possible **outcomes**
 - E.g., all the possible combinations of your DNA with the one of your {girl|boy}friend
 - You may or may not be able to observe directly the outcome.
 - \mathcal{F} if the set of **events** where an event is a set containing zero or more outcomes
 - E.g., the event of "the DNA corresponds to a girl with blue eyes"
 - An event is somehow more tangible and can generally be observed

Probabilities

- Using probabilities enables to model **uncertainty** that may result of **incomplete information** or **imprecise measurements**

A **random variable** (or stochastic variable) is, roughly speaking, a variable whose value results from a measurement (or an observation)

You can think of it as a **small box**:

- Every time you open the box, you get a different value.
 - I will use this box analogy throughout the whole lecture and I encourage you to ask yourself what the box can be in your own studies
- Formally a **probability space** is defined by (Ω, \mathcal{F}, P) where:
 - Ω , the **sample space**, is the set of all possible **outcomes**
 - E.g., all the possible combinations of your DNA with the one of your {girl|boy}friend
 - You may or may not be able to observe directly the outcome.
 - \mathcal{F} if the set of **events** where an event is a set containing zero or more outcomes
 - E.g., the event of "the DNA corresponds to a girl with blue eyes"
 - An event is somehow more tangible and can generally be observed
 - The **probability measure** $P : \mathcal{F} \rightarrow [0, 1]$ is a function returning an event's probability ($P(\text{"having a brown-eyed baby girl"}) = 0.0005$)

Continuous random variable

- A **random variable** associates a **numerical value** to **outcomes**

$$X : \Omega \rightarrow \mathbb{R}$$


- E.g., the weight of the baby at birth (assuming it solely depends on DNA, which is quite false but it's for the sake of the example)
- Since many computer science experiments are based on time measurements, we focus on **continuous** variables
- **Note:** To distinguish random variables, which are complex objects, from other mathematical objects, they will always be written in blue capital letters in this set of slides (e.g., X)
- The probability measure on Ω induces probabilities on the **values** of X
 - $P(X = 0.5213)$ is generally 0 as the outcome never exactly matches
 - $P(0.5213 \leq X \leq 0.5214)$ may however be non-zero

Probability distribution

A **probability distribution** (a.k.a. **probability density function** or p.d.f.) is used to describe the probabilities of different **values** occurring

- A random variable X has density f_X , where f_X is a non-negative and integrable function, if:

$$P[a \leq X \leq b] = \int_a^b f_X(w) dw$$




Note: the X in $1 \leq X \leq 6$ should be in blue...

- Note:** people often confuse the sample space with the random variable. Try to make the difference when modeling your system, it will help you

Characterizing a random variable

The probability density function **fully characterizes** the random variable but it is also complex object

- It may be symmetrical or not
- It may have one or several **modes**
- It may have a bounded support or not, hence the random variable may have a **minimal** and/or a **maximal** value
- The **median** cuts the probabilities in half



These are interesting aspects of f_X but they barely summarize it

Expected value and variance

- When one speaks of the "expected price", "expected height", etc. one means the **expected value** of a random variable that is a price, a height, etc.

$$E[X] = x_1 p_1 + x_2 p_2 + \dots + x_k p_k = \int_{-\infty}^{\infty} x f_X(x) dx$$

The expected value of X is the "average value" of X .

It is **not** the most probable value. The mean is one aspect of the distribution of X . The **median** or the **mode** are other interesting aspects.

- The **variance** is a measure of how far the values of a random variable are spread out from each other.

If a random variable X has the expected value (mean) $\mu = E[X]$, then the variance of X is given by:

$$\text{Var}(X) = E[(X - \mu)^2] = \int_{-\infty}^{\infty} (x - \mu)^2 f_X(x) dx$$

- The **standard deviation** σ is the square root of the variance. This normalization allows to compare it with the expected value

Outline

① A (mathematical) probabilistic model

② Using the model to estimate the expected value

Estimation

Evaluating and Comparing Alternatives With Confidence Intervals

What should I take care of?

③ Design of Experiments

Early Intuition and Key Concepts

④ Other random topics

Getting rid of Outliers

Summarizing the distribution

Estimating something else than the mean

Statistical Tests

References

How to estimate the Expected value?

To empirically **estimate** the expected value of a random variable X , one repeatedly measures observations of the variable and computes the arithmetic mean of the results

This is called the **sample mean**

Unfortunately, if you repeat the estimation, you may get a different value since X is a random variable ...

Central Limit Theorem [CLT]

- Let $\{X_1, X_2, \dots, X_n\}$ be a random sample of size n (i.e., a sequence of independent and identically distributed random variables with expected values μ and variances σ^2)
- The sample mean of these random variables is:

$$S_n = \frac{1}{n}(X_1 + \dots + X_n)$$


S_n is a random variable too!

- For large n 's, the distribution of S_n is approximately normal with mean μ and variance $\frac{\sigma^2}{n}$


$$S_n \xrightarrow{n \rightarrow \infty} \mathcal{N}\left(\mu, \frac{\sigma^2}{n}\right)$$

CLT Illustration: the mean smooths distributions

Start with an **arbitrary** distribution and compute the distribution of S_n for increasing values of n .




The Normal Distribution



The smaller the variance the more “spiky” the distribution.

The Normal Distribution




The smaller the variance the more “spiky” the distribution.

- Dark blue is less than one standard deviation from the mean. For the normal distribution, this accounts for about 68% of the set.
- Two standard deviations from the mean (medium and dark blue) account for about 95%
- Three standard deviations (light, medium, and dark blue) account for about 99.7%

Outline

- 1 A (mathematical) probabilistic model
- 2 Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- 3 Design of Experiments
 - Early Intuition and Key Concepts
- 4 Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References


CLT consequence: confidence interval



When n is large:

$$P\left(\mu \in \left[S_n - 2\frac{\sigma}{\sqrt{n}}, S_n + 2\frac{\sigma}{\sqrt{n}}\right]\right) = P\left(S_n \in \left[\mu - 2\frac{\sigma}{\sqrt{n}}, \mu + 2\frac{\sigma}{\sqrt{n}}\right]\right) \approx 95\%$$

CLT consequence: confidence interval




When n is large:

$$P\left(\mu \in \left[S_n - 2\frac{\sigma}{\sqrt{n}}, S_n + 2\frac{\sigma}{\sqrt{n}}\right]\right) = P\left(S_n \in \left[\mu - 2\frac{\sigma}{\sqrt{n}}, \mu + 2\frac{\sigma}{\sqrt{n}}\right]\right) \approx 95\%$$

There is 95% of chance that the **true mean** lies within $2\frac{\sigma}{\sqrt{n}}$ of the **sample mean**.

Without any particular hypothesis

- Assume, you have evaluated two **alternatives** A and B on n different **setups**
- You therefore consider the associated random variables A and B and try to estimate their expected values μ_A and μ_B




The two 95% confidence intervals do not overlap

$\leadsto \mu_A < \mu_B$ with more than 90% of confidence 😊

Without any particular hypothesis

- Assume, you have evaluated two **alternatives** A and B on n different **setups**
- You therefore consider the associated random variables A and B and try to estimate their expected values μ_A and μ_B




The two 95% confidence intervals do overlap

~> Nothing can be concluded 😞

Reduce C.I.?

Without any particular hypothesis

- Assume, you have evaluated two **alternatives** A and B on n different **setups**
- You therefore consider the associated random variables A and B and try to estimate their expected values μ_A and μ_B




The two 70% confidence intervals do not overlap

$\leadsto \mu_A < \mu_B$ with less than 50% of confidence 😞 \leadsto more experiments...

Without any particular hypothesis

- Assume, you have evaluated two **alternatives** A and B on n different **setups**
- You therefore consider the associated random variables A and B and try to estimate their expected values μ_A and μ_B




The width of the confidence interval is proportional to $\frac{\sigma}{\sqrt{n}}$

Halving C.I. requires 4 times more experiments! 😞

Try to **reduce variance** if you can... 😊


Exploiting blocks

- C.I.s overlap because variance is large. Some *setups* may have an intrinsically longer duration than others, hence a large $\text{Var}(A)$ and $\text{Var}(B)$



Exploiting blocks

- C.I.s overlap because variance is large. Some *setup*s may have an intrinsically longer duration than others, hence a large $\text{Var}(A)$ and $\text{Var}(B)$



- The previous test estimates μ_A and μ_B **independently**.


$$E[A] < E[B] \Leftrightarrow E[B - A] > 0.$$

In the previous evaluation, the **same** setup i is used for measuring A_i and B_i , hence we can focus on $B - A$.

Since $\text{Var}(B - A)$ is much smaller than $\text{Var}(A)$ and $\text{Var}(B)$, we can conclude that $\mu_A < \mu_B$ with 95% of confidence.


Exploiting blocks

- C.I.s overlap because variance is large. Some *setup*s may have an intrinsically longer duration than others, hence a large $\text{Var}(A)$ and $\text{Var}(B)$



- The previous test estimates μ_A and μ_B **independently**.
 $E[A] < E[B] \Leftrightarrow E[B - A] > 0$.
In the previous evaluation, the **same** setup i is used for measuring A_i and B_i , hence we can focus on $B - A$.
Since $\text{Var}(B - A)$ is much smaller than $\text{Var}(A)$ and $\text{Var}(B)$, we can conclude that $\mu_A < \mu_B$ with 95% of confidence.
- Relying on such common points is called **blocking** and enable to **reduce variance**.

Let's reuse a previous example



μ_A is 0.65 seconds smaller than μ_B with more than 99% of confidence 😊

You need to invest in a probabilistic model. Here we assumed:

- $A_i = \boxed{S_i} + A'_i$
- $B_i = \boxed{S_i} + B'_i$


So we could subtract them 😊

Dividing them would have been a very bad idea... 😞

How to compute and plot CI in R: code

```
1 library(ggplot2)
2 library(dplyr)
3 library(tidyr)
4 df = read.csv("data/set1.csv",header=T)
5 df$Diff=df$B-df$A # Assuming observations are paired!
6 dfgg = df %>% gather(Alternative, Time)
7 dfsum = dfgg %>%
8     group_by(Alternative) %>%
9     summarise(num = n(), mean = mean(Time), sd = sd(Time),
10     se = 2*sd/sqrt(num))
11 ggplot(dfgg,aes(x=Alternative,y=Time,color=Alternative)) +
12     scale_color_brewer(palette="Set1") + theme_bw() +
13     geom_jitter(alpha=.2,position = position_jitter(width = .1)) +
14     geom_errorbar(data=dfsum,width=.2,
15     aes(y=mean,ymin=mean-se,ymax=mean+se)) +
16     geom_point(data=dfsum,shape=21, size=3,
17     aes(y=mean,color=Alternative))
```

How to compute and plot CI in R: output




Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

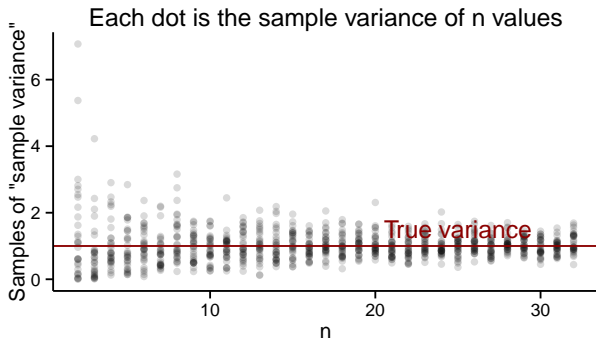
CLT hypothesis

- The CLT hypothesis are very weak: it **does not assume any particular distribution** (e.g., normality) for X
But, the CLT says that *when n goes large*, the sample mean is *normally distributed*. We have seen it holds true quickly



CLT hypothesis

- The CLT hypothesis are very weak: it **does not assume any particular distribution** (e.g., normality) for X
But, the CLT says that *when n goes large*, the sample mean is *normally distributed*. We have seen it holds true quickly
- However, the CLT uses $\sigma = \sqrt{\text{Var}(X)}$ but we only have the **sample variance**, not the **true variance**




CLT hypothesis

- The CLT hypothesis are very weak: it **does not assume any particular distribution** (e.g., normality) for X

But, the CLT says that *when n goes large*, the sample mean is *normally distributed*. We have seen it holds true quickly

- However, the CLT uses $\sigma = \sqrt{\text{Var}(X)}$ but we only have the **sample variance**, not the **true variance**

So you should always try to either find an **upper bound on the true variance** or **overestimate the sample variance** (e.g., `se=4*sd/sqrt(num)`)



How many replicates?

- **Q:** How Many Replicates?

How many replicates?

- **Q:** How Many Replicates?

A1: How many can you afford?

How many replicates?

- **Q:** How Many Replicates?

A1: How many can you afford?

A2: 30...

Rule of thumb: a sample of 30 or more is big sample but a sample of 30 or less is a small one (doesn't always work)

How many replicates?

- **Q:** How Many Replicates?

A1: How many can you afford?

A2: 30...

Rule of thumb: a sample of 30 or more is big sample but a sample of 30 or less is a small one (doesn't always work)

- With less than 30, you should make the C.I. wider using e.g., the **Student law**.

How many replicates?

- **Q:** How Many Replicates?

A1: How many can you afford?

A2: 30...

Rule of thumb: a sample of 30 or more is big sample but a sample of 30 or less is a small one (doesn't always work)

- With less than 30, you should make the C.I. wider using e.g., the **Student law**.
- Once you have a first C.I. with 30 samples, you can estimate how many samples will be required to answer your question. If it is too large, then either try to reduce variance (or the scope of your experiments) or simply explain that the two alternatives are hardly distinguishable... You need a **sequential approach**.

How many replicates?

- **Q:** How Many Replicates?

A1: How many can you afford?

A2: 30...

Rule of thumb: a sample of 30 or more is big sample but a sample of 30 or less is a small one (doesn't always work)

- With less than 30, you should make the C.I. wider using e.g., the **Student law**.
- Once you have a first C.I. with 30 samples, you can estimate how many samples will be required to answer your question. If it is too large, then either try to reduce variance (or the scope of your experiments) or simply explain that the two alternatives are hardly distinguishable... You need a **sequential approach**.
- **Running the right number of experiments enables to get to conclusions more quickly and hence to test other hypothesis.**





Key Hypothesis

The hypothesis of CLT are very weak. Yet, to qualify as replicates, the repeated measurements:

- must be **independent** (take care of warm-up)
- must **not** be part of a **time series** (the system behavior may temporary change)
- must **not** come **from the same place** (the machine may have a problem)
- must be of appropriate **spatial scale**

Perform graphical checks

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis





Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear

If you see several modes, there may be an hidden parameter to take into account

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis





Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear

If you see several modes, there may be an hidden parameter to take into account

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis





Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear

If you see several modes, there may be an hidden parameter to take into account

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis





Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear

If you see several modes, there may be an hidden parameter to take into account

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis





Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear

If you see several modes, there may be an hidden parameter to take into account

Simple Graphical Checks



Fixed Location the run sequence plot should be flat and non-drifting

Fixed Variation the vertical spread in the run sequence plot should approximately the same over the entire horizontal axis



Independence the lag plot should be structureless

Fixed Distribution (, in particular if the *fixed normal distribution* assumption holds)

- the histogram should be bell-shaped, and
- the normal probability plot should be linear


If you see several modes, there may be an hidden parameter to take into account

Temporal Dependency







- Should look independent and statistically identical
- **Periodicity** : May depend on sampling frequency or on clock resolution
 - Study the period (Fourier), use time series
- **Danger**: temporal correlation \leadsto study **stationarity**

Detect Trends



- Model the trend: here increases then saturates
- Possibly remove the trend by compensating it (multiplicative factor here) or removing what can be identified as a warm-up



Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Comparing Two Alternatives (Blocking + Randomization)

- When comparing A and B for different settings, doing A, A, A, A, A, A and then B, B, B, B, B, B is a bad idea

Comparing Two Alternatives (Blocking + Randomization)

- When comparing A and B for different settings, doing A, A, A, A, A, A and then B, B, B, B, B, B is a bad idea
- You should better do $A, B, \quad A, B, \quad A, B, \quad A, B, \dots$

Comparing Two Alternatives (Blocking + Randomization)

- When comparing A and B for different settings, doing A, A, A, A, A, A and then B, B, B, B, B, B is a bad idea
- You should better do $A, B, A, B, A, B, A, B, \dots$
- Even better, randomize your run order. You should flip a coin for each configuration and start with A on head and with B on tail. . .

$A, B, B, A, B, A, A, B, \dots$

With such design, you will even be able to check whether being the first alternative to run changes something or not

Comparing Two Alternatives (Blocking + Randomization)

- When comparing A and B for different settings, doing A, A, A, A, A, A and then B, B, B, B, B, B is a bad idea
- You should better do A, B, A, B, A, B, A, B, . . .
- Even better, randomize your run order. You should flip a coin for each configuration and start with A on head and with B on tail. . .

A, B, B, A, B, A, A, B, . . .

With such design, you will even be able to check whether being the first alternative to run changes something or not

- Each configuration you test should be run on different machines
You should record as much information as you can on how the experiments was performed

Experimental Design

There are two key concepts:

replication and randomization

You replicate to increase reliability. You randomize to reduce bias.

**If you replicate thoroughly and randomize properly,
you will not go far wrong.**

Experimental Design

There are two key concepts:

replication and **randomization**

You replicate to **increase reliability**. You randomize to **reduce bias**.

**If you replicate thoroughly and randomize properly,
you will not go far wrong.**

It doesn't matter if you cannot do your own advanced statistical analysis. If you designed your experiments properly, you may be able to find somebody to help you with the statistics.

If your experiments is not properly designed, then no matter how good you are at statistics, you experimental effort will have been wasted.

No amount of high-powered statistical analysis can turn a bad experiment into a good one.

Other important concepts:

- **Pseudo-replication**
- **Experimental vs. observational data**

Replication vs. Pseudo-replication

Measuring the same configuration several times is not replication. It's **pseudo-replication** and is generally biased

Instead, test **other** configurations (with a good randomization)

In case of pseudo-replication, here is what you can do:

- average away the pseudo-replication and carry out your statistical analysis on the means
- carry out separate analysis for each time period
- use proper time series analysis

Experimental data vs. Observational data

You need a good blend of **observation**, **theory** and **experiments**

- Many scientific experiments appear to be carried out with no hypothesis in mind at all, but simply to see what happens.
- This may be OK in the early stages but drawing conclusions on such observations is difficult (large number of equally plausible explanations; without testable prediction no experimental ingenuity; ...).

Experimental data vs. Observational data

You need a good blend of **observation**, **theory** and **experiments**

- Many scientific experiments appear to be carried out with no hypothesis in mind at all, but simply to see what happens.
- This may be OK in the early stages but drawing conclusions on such observations is difficult (large number of equally plausible explanations; without testable prediction no experimental ingenuity; ...).

Strong inference Essential steps:

- ① Formulate a clear hypothesis
- ② Devise an acceptable test

Weak inference It would be silly to disregard all observational data that do not come from designed experiments. Often, they are the only we have (e.g. the trace of a system).

But we need to keep the limitations of such data in mind. It is possible to use it to **derive hypothesis** but not to **test hypothesis** (i.e., **claim facts**).

Correlation and Causation

Let me illustrate this inference story with a few examples.

It may be the case that two random variables X and Y are **dependent**

- E.g., Let's pick a student at random and measure its *TimeSpentStudying* and its *TestScore*
 - In most cases, studying more should improve your test score 😊


The **correlation** of two variables X and Y is defined as:

$$\text{corr}(X, Y) = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y} = \frac{E[(X - \mu_X)(Y - \mu_Y)]}{\sigma_X \sigma_Y}$$

- The correlation is symmetrical ($\text{corr}(X, Y) = \text{corr}(Y, X)$)
- The correlation is in $[-1, 1]$
- $\text{corr}(Y, X) = 1$ or $-1 \Rightarrow$ perfectly linear relationship
- X independent of $Y \Rightarrow \text{corr}(X, Y) = 0$
- Y grows when X grows $\Rightarrow \text{corr}(X, Y) > 0$

It is thus very tempting to use **sample correlation** as a way of knowing whether some variables are **dependant**



Correlation does not imply Causation



Mikhail Ryazanov (talk) - PiratesVsTemp.svg.
Licensed under CC BY-SA 3.0 via Wikimedia Commons

- 2 variables peuvent être fortement corrélées à une troisième (e.g., year)
- Btw, what is wrong with this figure? 😊

Observational vs. Experimental Data Illustration




Source: *Spurious correlations*. For the good of the US society, we should try to get rid of honey bees 😊

Correlation does not imply Causation

For any two correlated events, A and B, the following relationships are possible:

- A causes B (direct causation) 😊
- A causes B and B causes A (bidirectional or cyclic causation) 😊
- A causes C which causes B (indirect causation) 😊
- B causes A; (reverse causation) 😞
- A and B are consequences of a common cause, but do not cause each other 😞
- There is no connection between A and B; it is a coincidence 😞
 - But **designed experiments** can help you ruling this option out




Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References


Abnormal measurements

- Rare events: interpretation
- Get rid of it using e.g., quantiles:
 - What is the good rejection rate? E.g., "above $Q3 + 1.5 \times (IQR)$ " (boxplot, Tukey, 1977), i.e., above $\mu + 2\sigma$ for a normal distribution?
- A threshold value: what is the right threshold?
 - Reject values larger than 100 \leadsto .6% of rejection
 - Reject values larger than 50 \leadsto 1% of rejection
 - Reject values larger than 10 \leadsto 6% of rejection




Abnormal measurements

- **Rare events:** interpretation
- Get rid of it using e.g., **quantiles**:
 - What is the good **rejection rate**? E.g., "above $Q3 + 1.5 \times (IQR)$ " (boxplot, Tukey, 1977), i.e., above $\mu + 2\sigma$ for a normal distribution?
- A threshold value: what is the right threshold?
 - Reject values larger than 100 \leadsto .6% of rejection
 - Reject values larger than 50 \leadsto 1% of rejection
 - Reject values larger than 10 \leadsto 6% of rejection



Abnormal measurements

- **Rare events:** interpretation
- Get rid of it using e.g., **quantiles**:
 - What is the good **rejection rate**? E.g., "above $Q3 + 1.5 \times (IQR)$ " (boxplot, Tukey, 1977), i.e., above $\mu + 2\sigma$ for a normal distribution?
- A threshold value: what is the right threshold?
 - Reject values larger than 100 \leadsto .6% of rejection
 - Reject values larger than 50 \leadsto 1% of rejection
 - Reject values larger than 10 \leadsto 6% of rejection






Actually, I generated these samples using the Cauchy distribution, which is pathological for most idea you'll come up with 😊

There is **no mathematical definition of what constitutes an outlier**. It's related to the experimenter's interpretation and is subjective...

Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Summarizing the distribution






What is the shape of the histogram:

- Uni/multi-modal:
 - If uni-modal, summarize with **central tendency** (mean, mode, median)
 - Symmetrical or not (\leadsto skewness)
 - Flat or not (\leadsto kurtosis)

If uni-modal you can go for a boxplot but avoid other fancy plots unless you know what you do...

Summarizing the distribution



What is the shape of the histogram:

- Uni/multi-modal:
 - If uni-modal, summarize with **central tendency** (mean, mode, median)
 - Symmetrical or not (\leadsto skewness)
 - Flat or not (\leadsto kurtosis)

If uni-modal you can go for a boxplot but avoid other fancy plots unless you know what you do...

Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Biased and unbiased estimators...

Expected value the *sample mean* is unbiased but is "sensitive" to outliers. This is not an excuse for estimating something else! Furthermore, there is an easy way to compute confidence intervals.

Mode the *naive estimate* is unstable and depends on the histogram's bin width. Still ongoing research on this:

E.g., Bickela and Frühwirthb, *On a fast, robust estimator of the mode: Comparisons to other robust estimators with applications*, Computational Statistics & Data Analysis 2006

Median the *sample median* is robust to outliers but it is quite sensitive to discrete distributions... There exists other more involved estimators but is median really what you want to estimate?

Minimum and Maximum the *sample minimum* is always too large, hence it is biased...

Variance `var` is a unbiased estimator of variance but `sd` is a biased estimator of standard deviation. Unbiasing depends on the distribution so just overestimate...

Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Tests

We have seen how to build **estimators** of specific characteristics of f_X based on observations of X

- Having an estimator is worth only if you can provide **confidence** on this estimation

Estimates can be used to **test** hypothesis

- We have seen how we could test whether $\mu_A < \mu_B$ from estimates of μ_A and μ_B

But there may be other more efficient ways to test such hypothesis

- if you know observations are paired
- if you know something about the underlying distribution

Other kind of complex hypothesis may tested

- $\text{median}(A) = \text{median}(B)$
- A and B follow the same distribution

We could give a whole lecture on this topic... Only use the tests you truly understand

Outline

- ① A (mathematical) probabilistic model
- ② Using the model to estimate the expected value
 - Estimation
 - Evaluating and Comparing Alternatives With Confidence Intervals
 - What should I take care of?
- ③ Design of Experiments
 - Early Intuition and Key Concepts
- ④ Other random topics
 - Getting rid of Outliers
 - Summarizing the distribution
 - Estimating something else than the mean
 - Statistical Tests
 - References

Roadmap for a good data analysis (Jain)

- 1 Plot the sample (various representations)
- 2 Describe the results (data analysis)
- 3 Preliminary processing : remove or flag outliers, estimate or flag missing values
- 4 Propose a stochastic model. Establish the hypothesis: independence (time correlation, auto-correlation), stationarity, same probability law
- 5 Summarize data by a histogram
- 6 Comment the shape (modal/skewness/flatness/...)
- 7 Estimate the central tendency of the sample : choose the central index
- 8 Estimate the accuracy of the result (confidence intervals)
- 9 Propose a visualization