

# OBTAINING FAITHFUL/REPRODUCIBLE MEASUREMENTS ON MODERN CPUS

---

Tom Cornebize and Arnaud Legrand

Univ. Grenoble Alpes, CNRS, Inria, Grenoble INP

Nantes, Journées du GDR RSD, January 2020

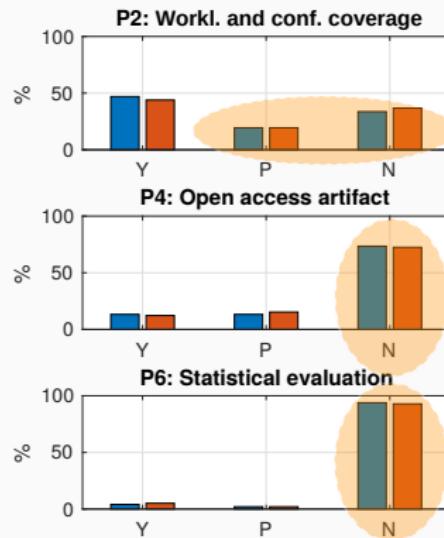
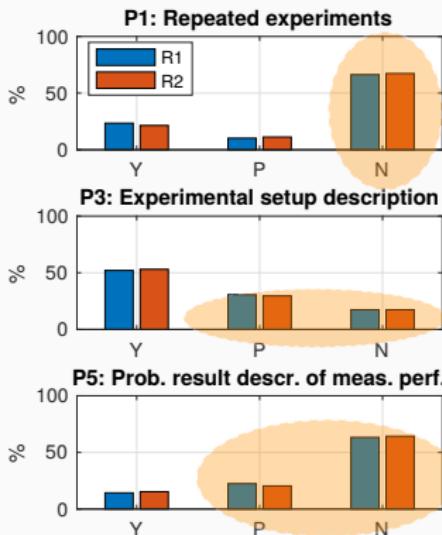


# TOWARD REPRODUCIBLE COMPUTER SCIENCE EXPERIMENTS ?

---

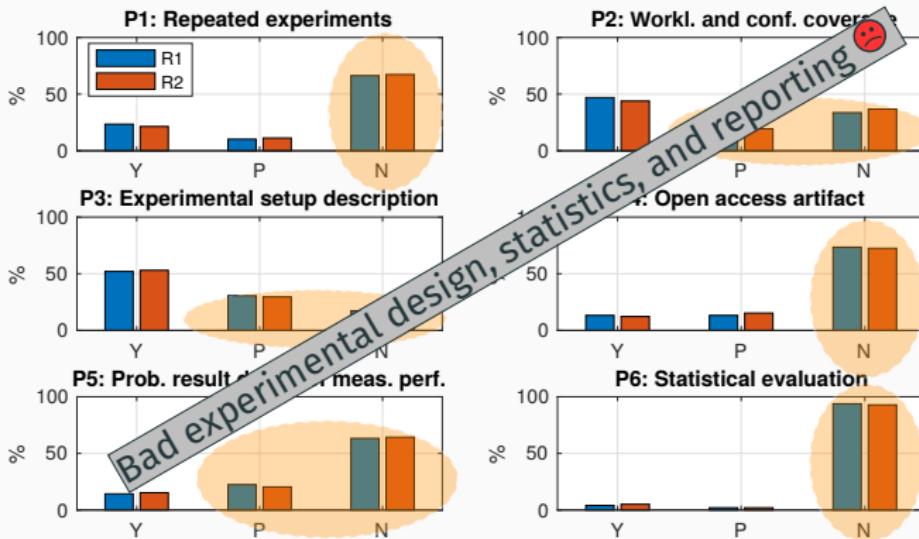
# KEY CONCERN FOR OUR COMMUNITY (ROOM FOR IMPROVEMENT)

How are cloud performance currently obtained and reported?, Methodological Principles for Reproducible Performance Evaluation in Cloud Computing, IEEE Trans. on Soft. Eng., July 2019



# KEY CONCERN FOR OUR COMMUNITY (ROOM FOR IMPROVEMENT)

How are cloud performance currently obtained and reported?, Methodological Principles for Reproducible Performance Evaluation in Cloud Computing, IEEE Trans. on Soft. Eng., July 2019



Key DoE principles:

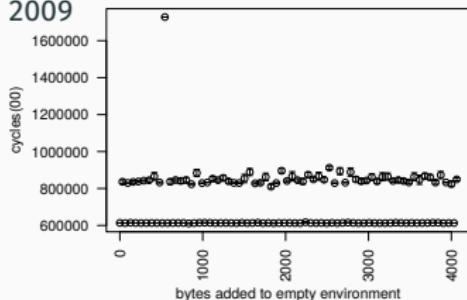
1. Replicate to increase reliability.
2. Randomize to reduce bias ↪ Evaluate statistical confidence.

# MEASURING PERFORMANCE IS DIFFICULT

*Producing wrong data without doing anything obviously wrong!*

Mytkowicz et al. in ACM SIGPLAN Not. 44(3), March 2009

*changing the size of environment variables can trigger performance degradation as high as 300%; simply changing the link order of object files can cause performance to decrease by as much as 57%.*

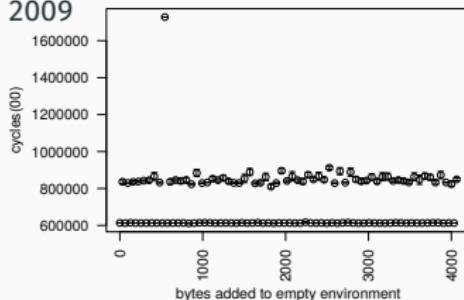


# MEASURING PERFORMANCE IS DIFFICULT

Producing wrong data without doing anything obviously wrong!

Mytkowicz et al. in ACM SIGPLAN Not. 44(3), March 2009

changing the size of *environment variables* can trigger performance degradation as high as 300%; simply changing the *link order* of object files can cause performance to decrease by as much as 57%.



Taming the Influence of Memory Layout. STABILIZER: Statistically Sound Performance Evaluation, C. Curtsinger and E. Berger in ASPLOS 2013

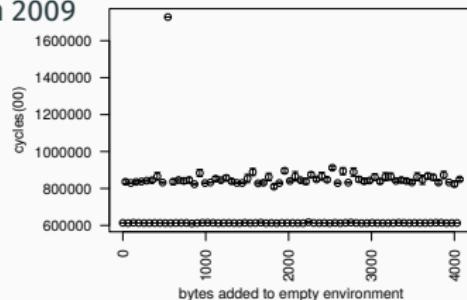
STABILIZER forces executions to sample the space of memory configurations by *repeatedly rerandomizing* layouts of code, stack, and heap objects at run-time. [...] Re-randomization ensures that layout effects *follow a Gaussian distribution*, enabling the use of statistical tests like ANOVA.

# MEASURING PERFORMANCE IS DIFFICULT

Producing wrong data without doing anything obviously wrong!

Mytkowicz et al. in ACM SIGPLAN Not. 44(3), March 2009

changing the size of *environment variables* can trigger performance degradation as high as 300%; simply changing the *link order* of object files can cause performance to decrease by as much as 57%.



Taming the Influence of Memory Layout. *STABILIZER: Statistically Sound Performance Evaluation*, C. Curtsinger and E. Berger in ASPLOS 2013

STABILIZER forces executions to sample the space of memory configurations by *repeatedly rerandomizing* layouts of code, stack, and heap objects at run-time. [...] Re-randomization ensures that layout effects *follow a Gaussian distribution*, enabling the use of statistical tests like ANOVA.

Randomization helps fighting bias incurred by:

- specific configurations      $AA \dots A \rightarrow A_1 A_2 \dots A_n$  (*pseudo-replication*)
- temporary perturbations    $AA \dots A BB \dots B \rightarrow ABBAAB \dots$

## CACHE EFFECTS!

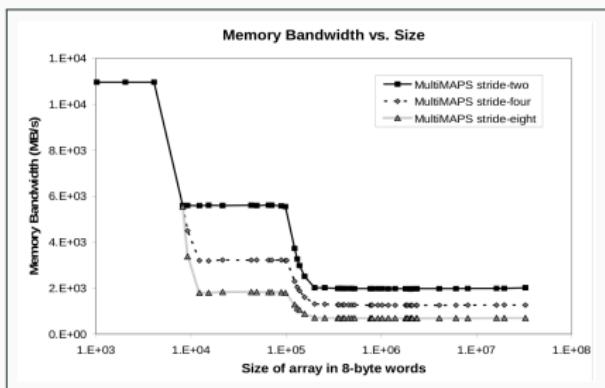
---

## IMPACT OF WORKING SET SIZE ON EFFECTIVE BANDWIDTH

- Cache hierarchy (L1, L2, L3, RAM) with different bandwidth
- LRU, pre-fetching for linear access
- An array fits in a cache level  $\Rightarrow$  operate at the corresponding bandwidth
- Stride access decrease bandwidth

# IMPACT OF WORKING SET SIZE ON EFFECTIVE BANDWIDTH

- Cache hierarchy (L1, L2, L3, RAM) with different bandwidth
- LRU, pre-fetching for linear access
- An array fits in a cache level  $\Rightarrow$  operate at the corresponding bandwidth
- Stride access decrease bandwidth



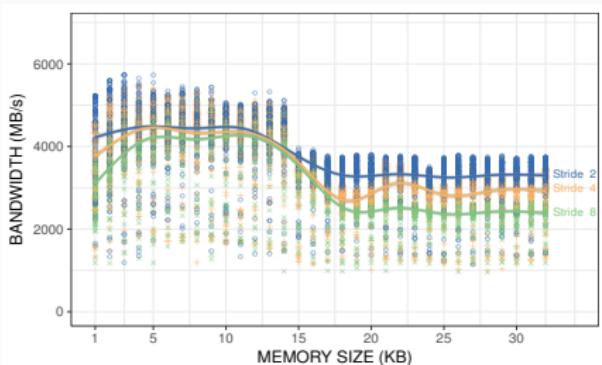
MultiMAPS on an Opteron

*Genetic Algorithms Approach to Modeling the Performance of Memory-bound Computations,*  
Tikir et. al. in SC'07

```
MultiMAPS(size, stride, nloops) {  
    allocate buffer[size];  
    timer_start();  
    for rep in (1..nloops)  
        for i in (0..size/stride)  
            access buffer[stride*i];  
    timer_stop();  
    bandwidth = nb_access /  
               elapsed_time;  
    deallocate buffer;  
}
```

# IMPACT OF WORKING SET SIZE ON EFFECTIVE BANDWIDTH

- Cache hierarchy (L1, L2, L3, RAM) with different bandwidth
- LRU, pre-fetching for linear access
- An array fits in a cache level  $\Rightarrow$  operate at the corresponding bandwidth
- Stride access decrease bandwidth

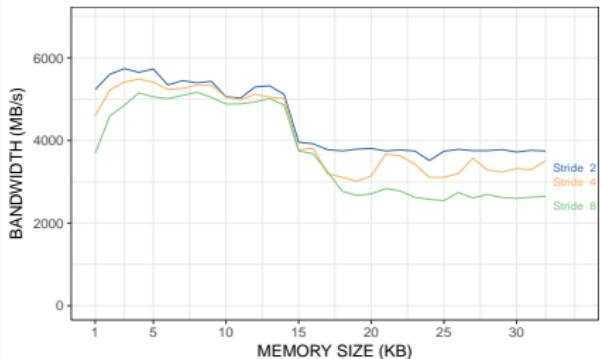


Our first attempt on a Pentium 4...

```
MultIMAPS(size, stride, nloops) {  
    allocate buffer[size];  
    timer_start();  
    for rep in (1..nloops)  
        for i in (0..size/stride)  
            access buffer[stride*i];  
    timer_stop();  
    bandwidth = nb_access /  
               elapsed_time;  
    deallocate buffer;  
}
```

# IMPACT OF WORKING SET SIZE ON EFFECTIVE BANDWIDTH

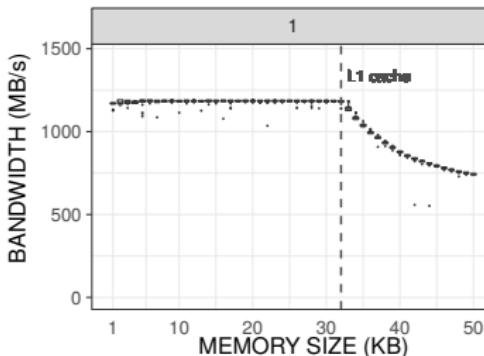
- Cache hierarchy (L1, L2, L3, RAM) with different bandwidth
- LRU, pre-fetching for linear access
- An array fits in a cache level  $\Rightarrow$  operate at the corresponding bandwidth
- Stride access decrease bandwidth



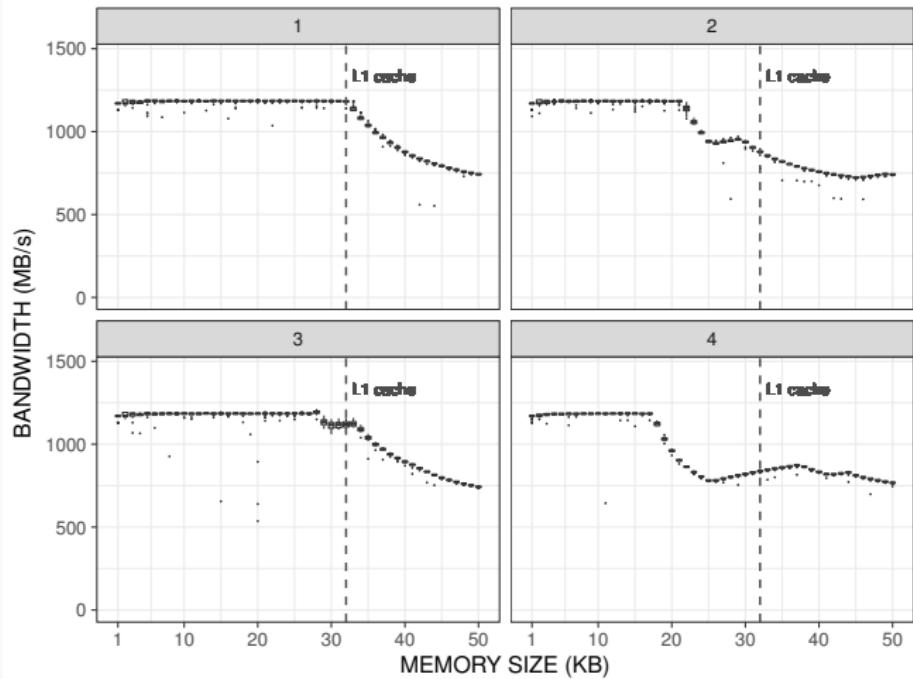
Our first attempt on a Pentium 4...

```
MultIMAPS(size, stride, nloops) {  
    allocate buffer[size];  
    timer_start();  
    for rep in (1..nloops)  
        for i in (0..size/stride)  
            access buffer[stride*i];  
    timer_stop();  
    bandwidth = nb_access /  
               elapsed_time;  
    deallocate buffer;  
}
```

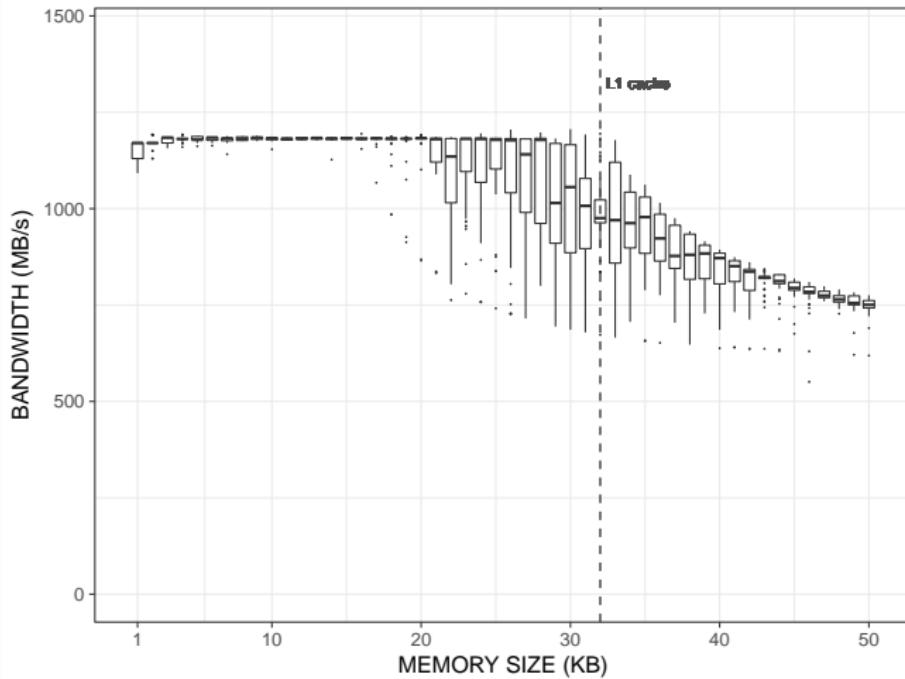
# IMPACT OF ARCHITECTURE



# IMPACT OF ARCHITECTURE



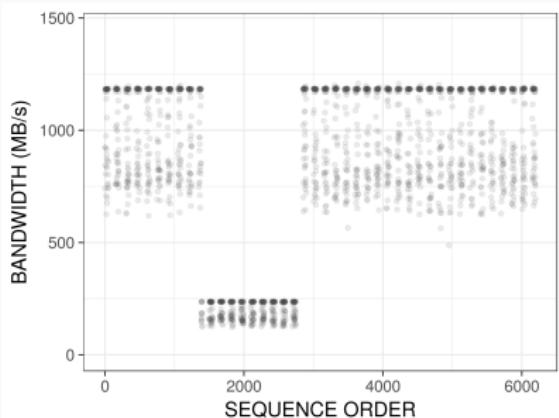
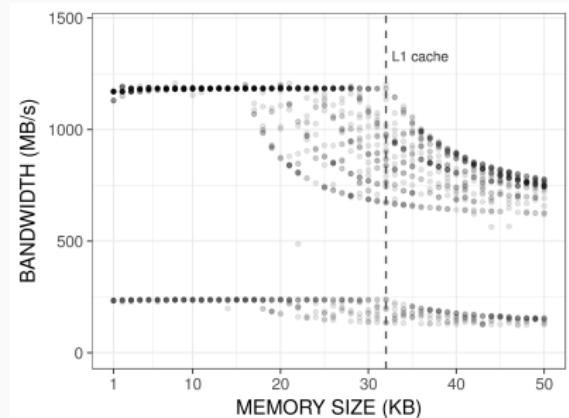
# IMPACT OF ARCHITECTURE (THE ARM ASSOCIATIVITY ISSUE)



Randomize physical address start!

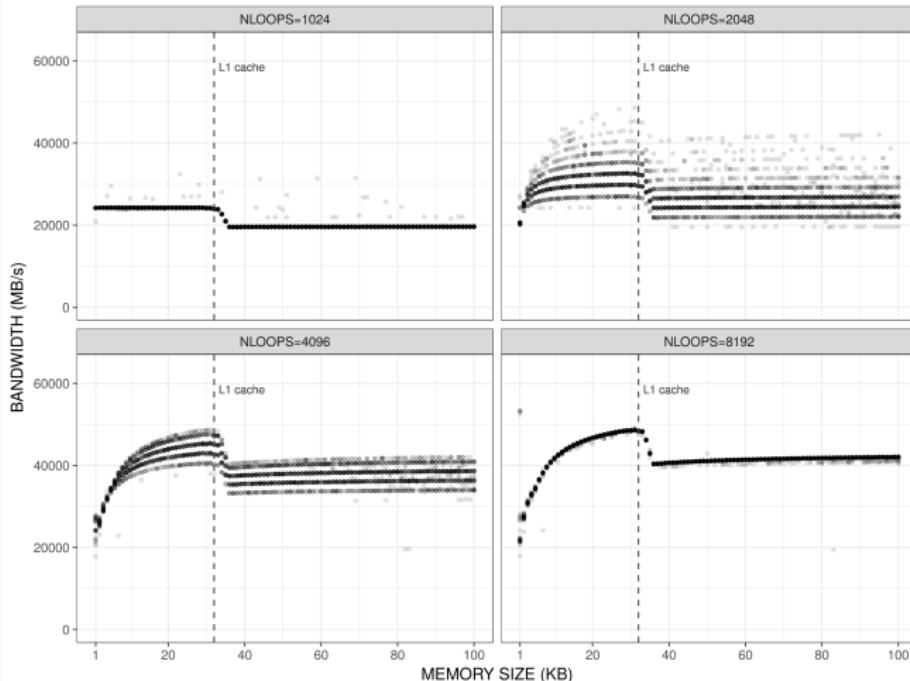
# IMPACT OF OPERATING SYSTEM SCHEDULER

- Activating real-time kernel scheduler



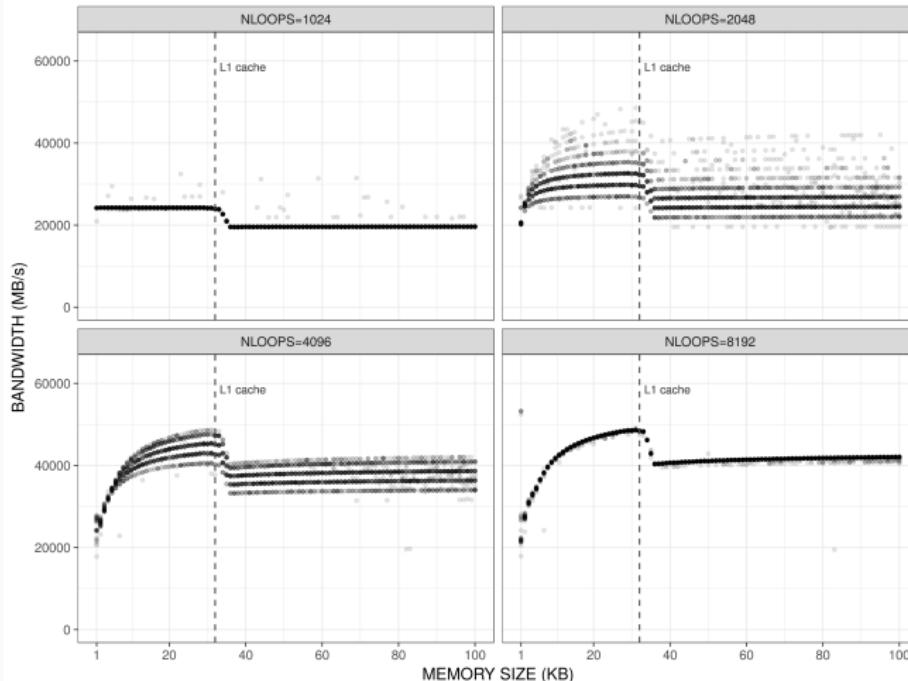
# IMPACT OF REPETITIONS

- Remember `nloops` ?



# IMPACT OF REPETITIONS DVFS

- Remember nloops ?

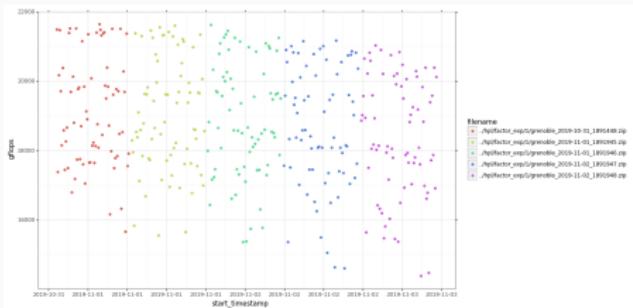


## CPU PERFORMANCE

---

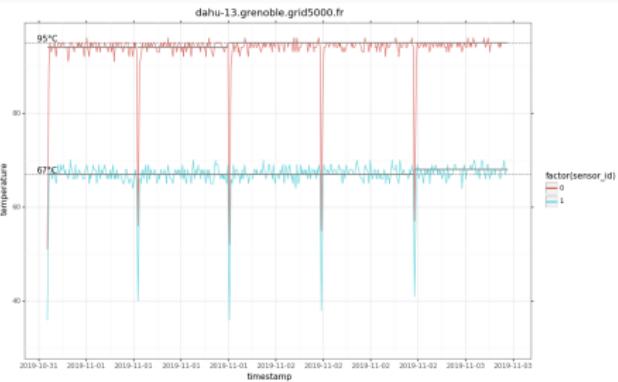
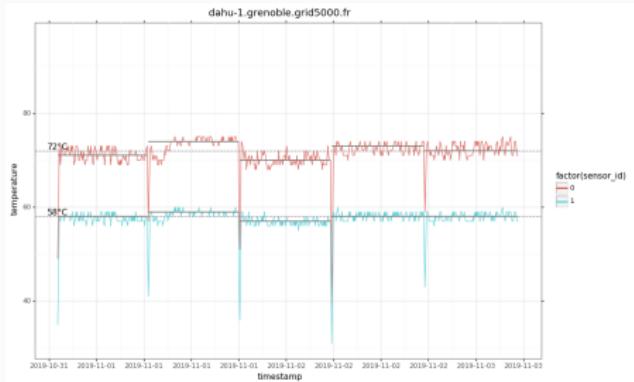
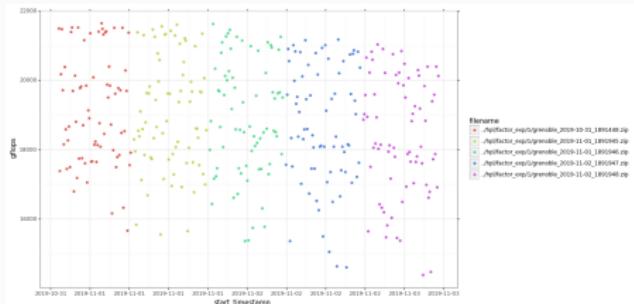
# AVOIDING "TEMPORARY" PERTURBATIONS (RANDOMIZING A FACTORIAL DESIGN)

- HPL performance (70 config., 5 repetitions)
- Time scale = 3 days



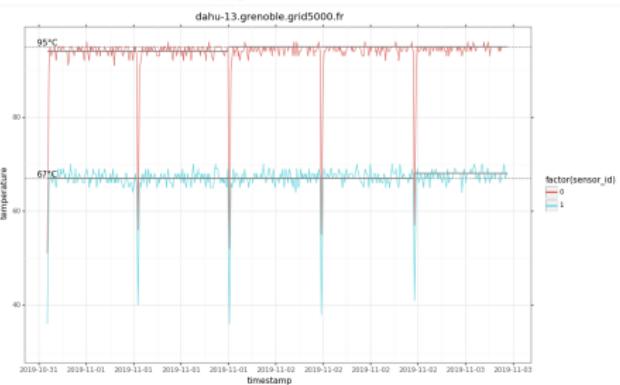
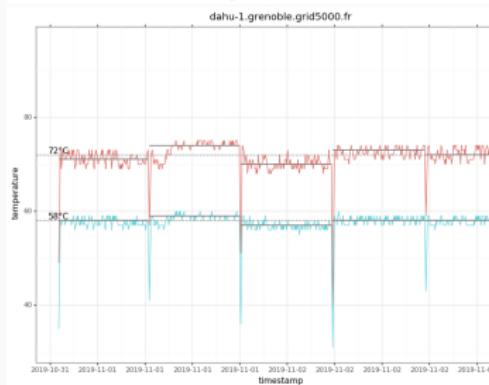
# AVOIDING "TEMPORARY" PERTURBATIONS (RANDOMIZING A FACTORIAL DESIGN)

- HPL performance (70 config., 5 repetitions)
- Time scale = 3 days



# AVOIDING "TEMPORARY" PERTURBATIONS (RANDOMIZING A FACTORIAL DESIGN)

- HPL performance (70 config., 5 repetitions)
- Time scale = 3 days



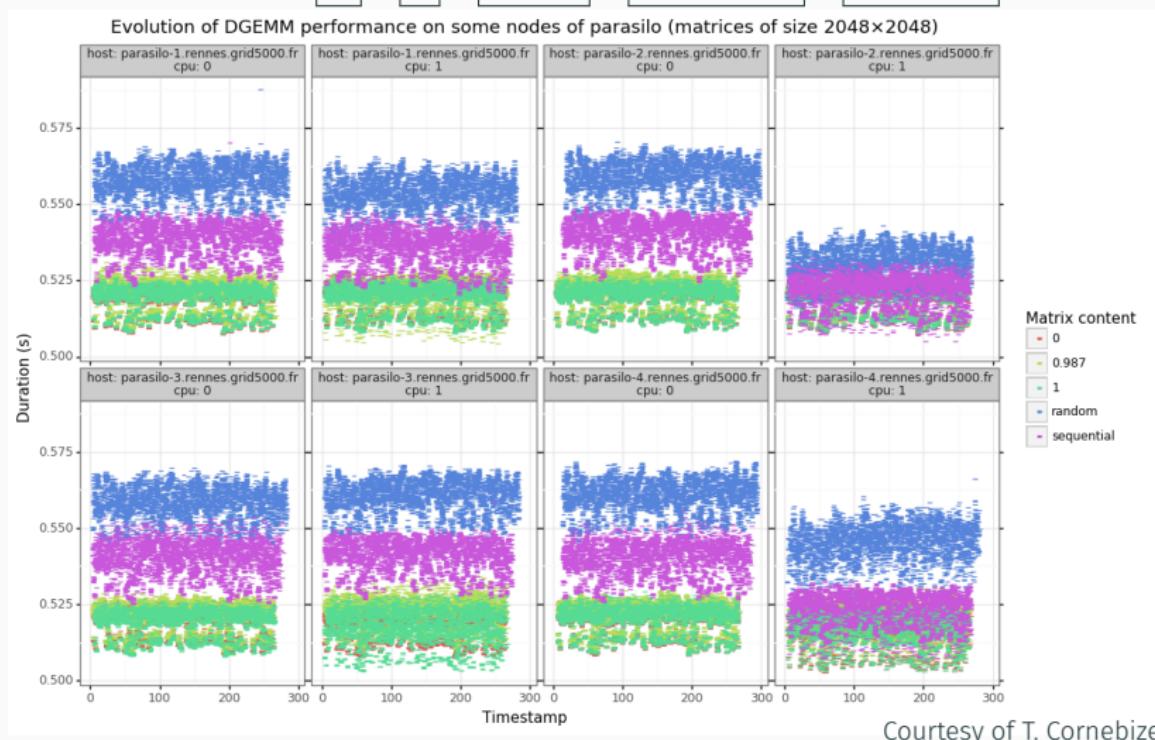
# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ ,  $2048 \times 2048$  matrix
- A initialized with
- Time scale = 5 minutes

# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ , 2048 × 2048 matrix
- A initialized with 

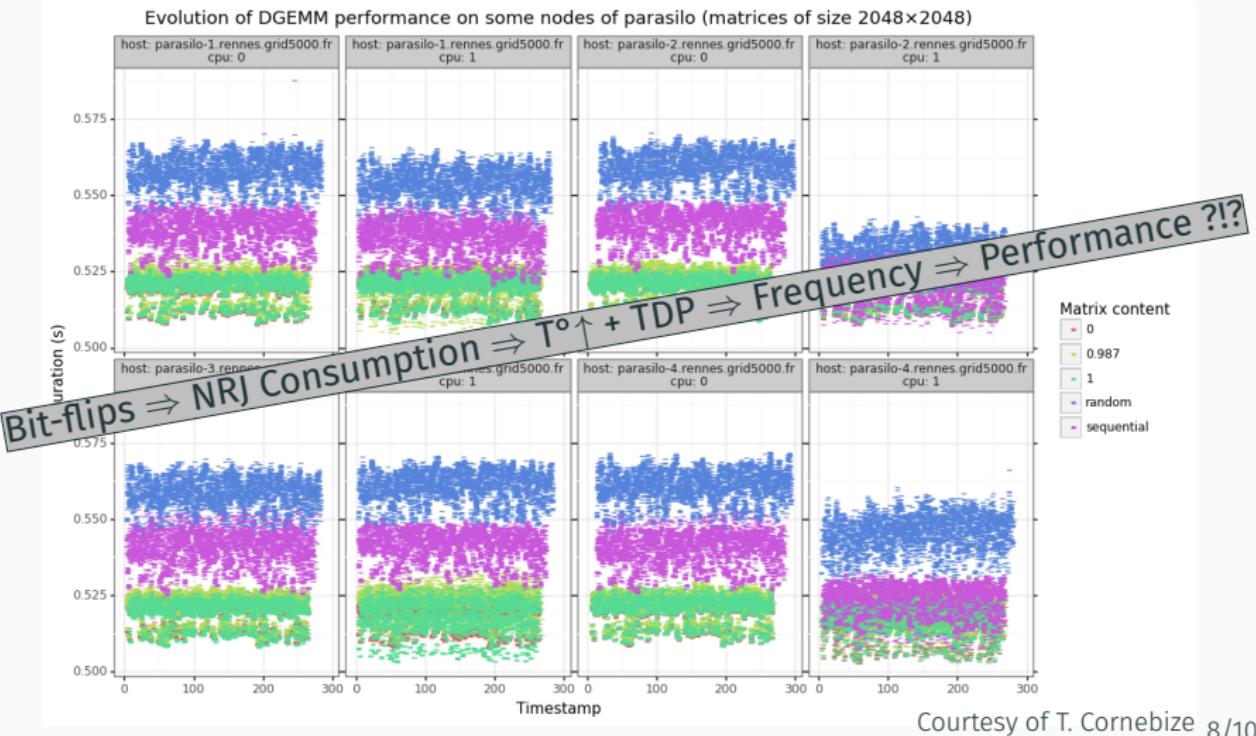
0	1	0.987	1, 2, 3, ...	random ?
---	---	-------	--------------	----------
- Time scale = 5 minutes



# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ , 2048 × 2048 matrix
- A initialized with 

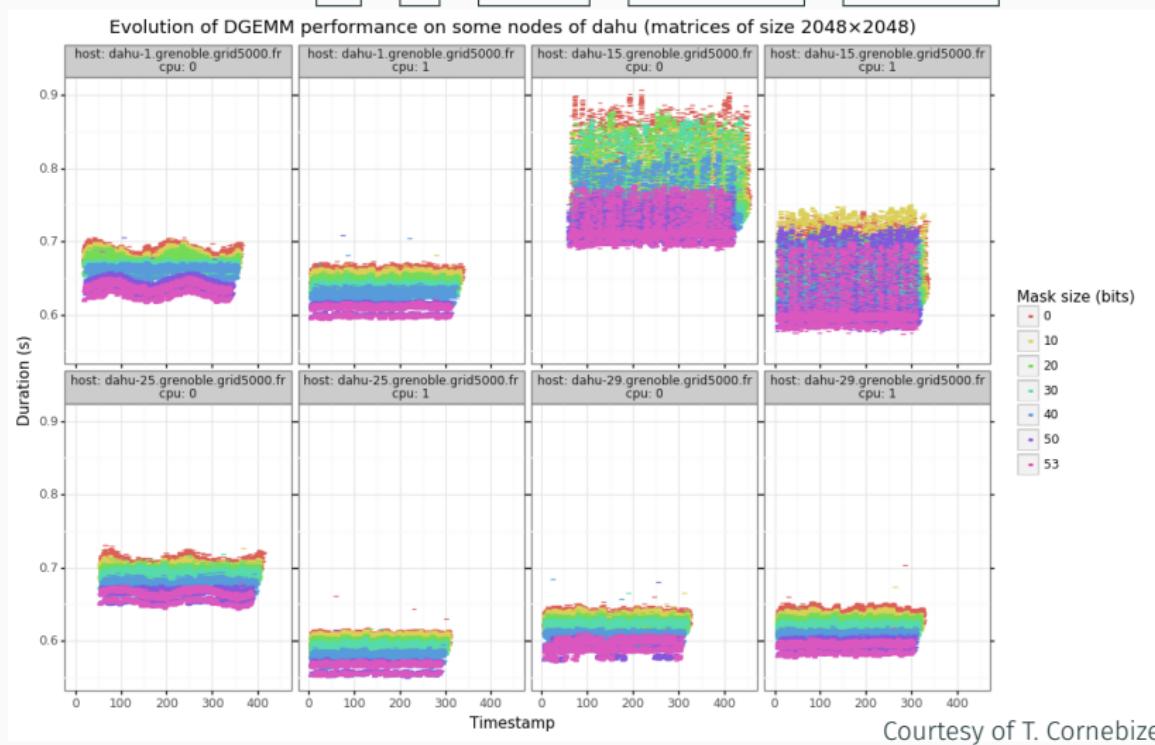
0	1	0.987	1, 2, 3, ...	random ?
---	---	-------	--------------	----------
- Time scale = 5 minutes



# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ ,  $2048 \times 2048$  matrix
- $A$  initialized with 

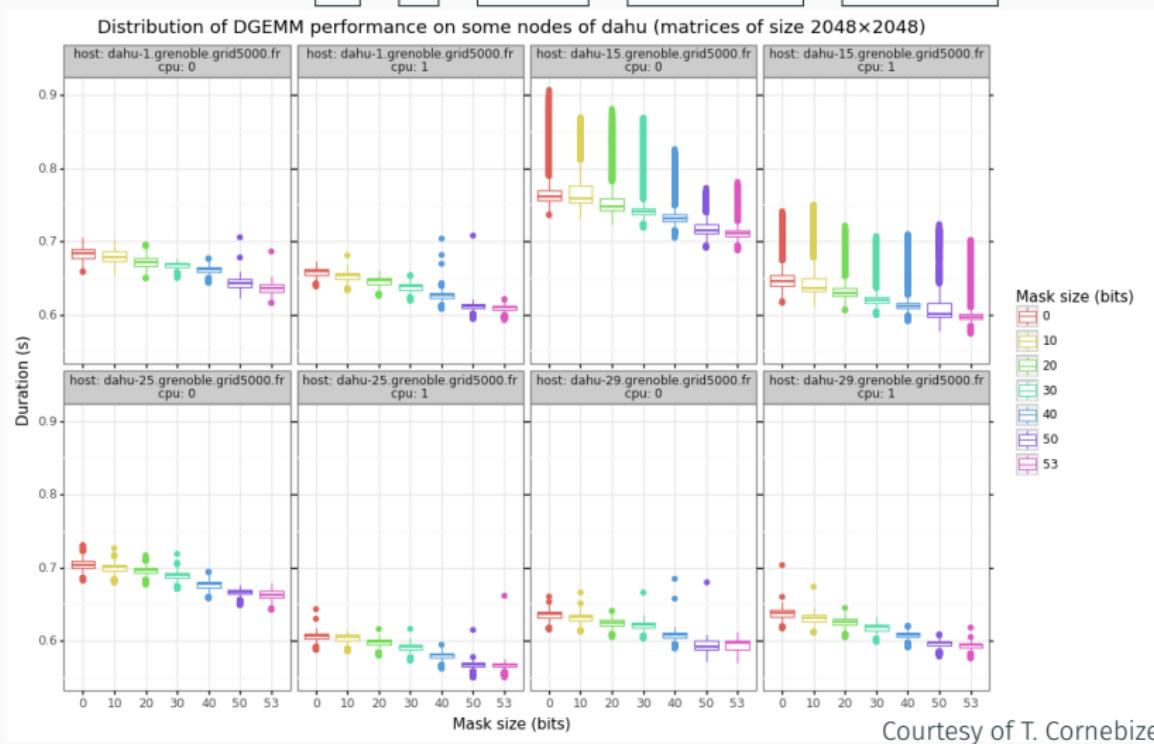
0	1	0.987	1, 2, 3, ...	random ?
---	---	-------	--------------	----------
- Time scale = 5 minutes



# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ , 2048 × 2048 matrix
- A initialized with 

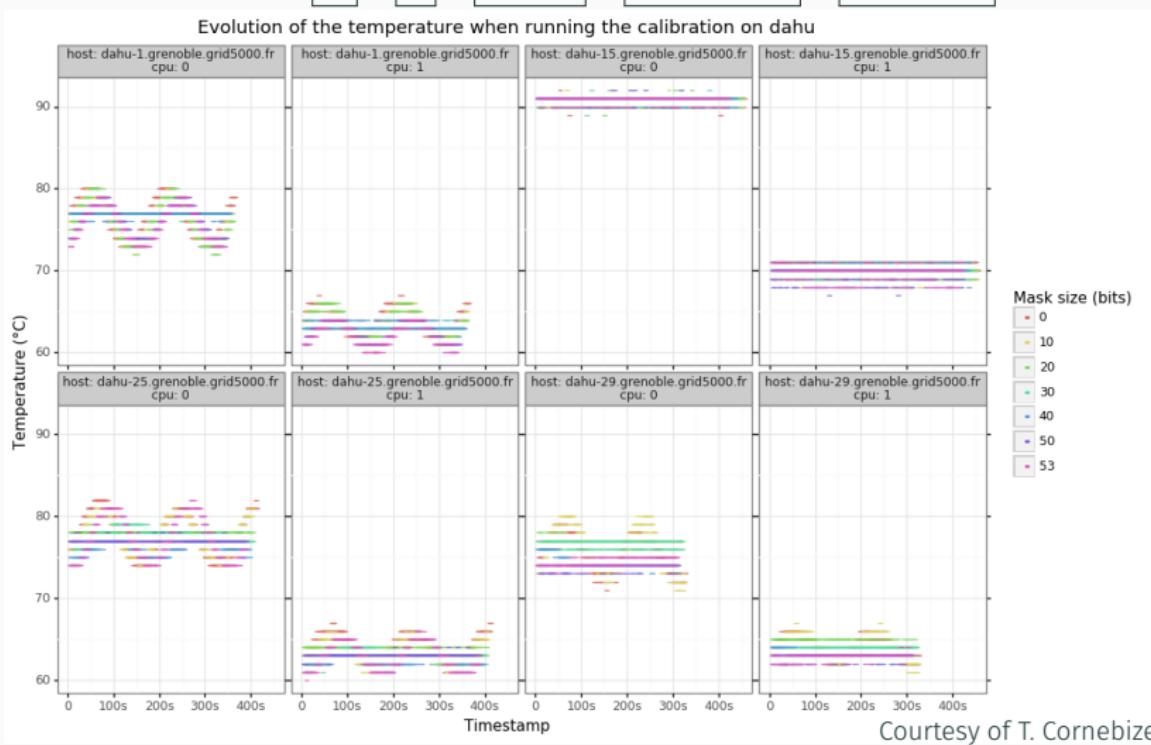
0	1	0.987
1, 2, 3, ...	random ?	
- Time scale = 5 minutes



# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ ,  $2048 \times 2048$  matrix
- $A$  initialized with 

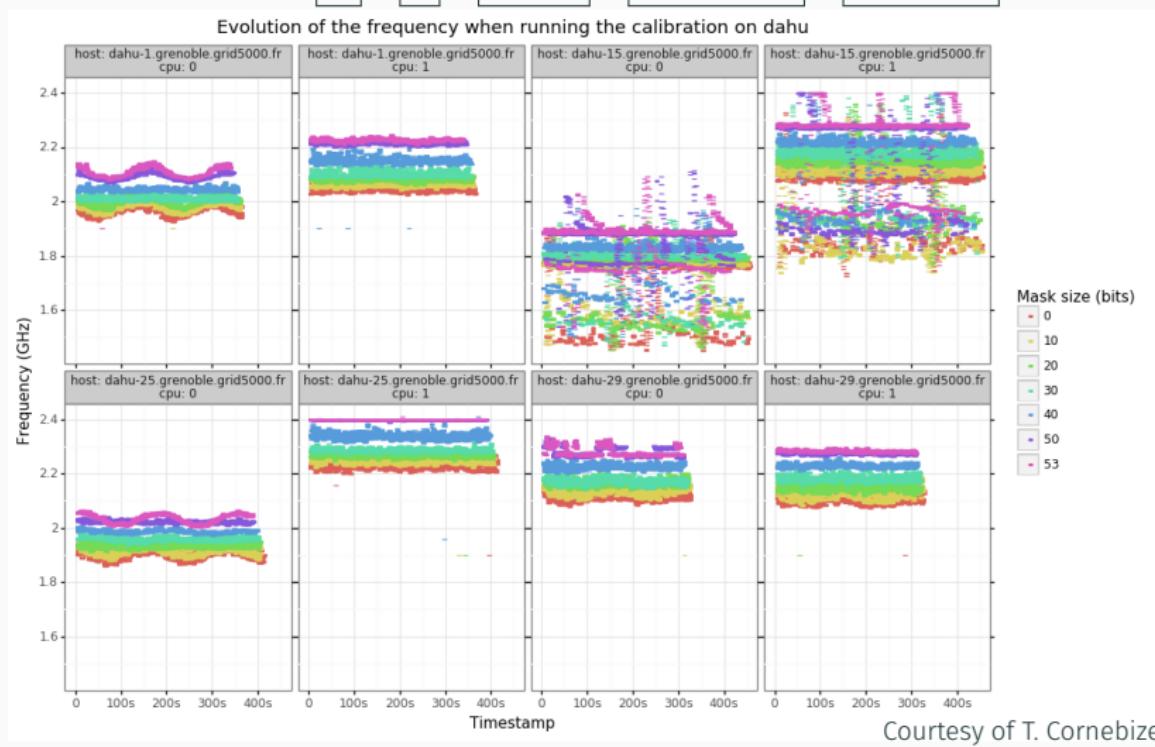
0	1	0.987	1, 2, 3, ...	random ?
---	---	-------	--------------	----------
- Time scale = 5 minutes



# ON THE IMPORTANCE OF CONTENT INITIALIZATION

- $C = A \times A$ ,  $2048 \times 2048$  matrix
- $A$  initialized with 

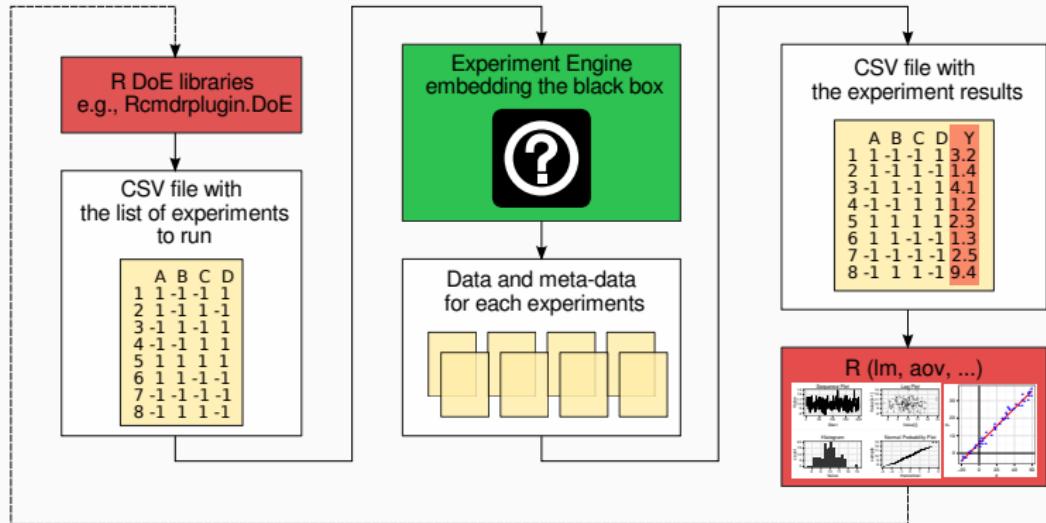
0	1	0.987	1, 2, 3, ...	random ?
---	---	-------	--------------	----------
- Time scale = 5 minutes



## CONCLUSION

---

# EXPERIMENTAL METHODOLOGY: NOTICING THE UNEXPECTED



1. A **separation of concerns**
  - Transparent Measurement Procedure and Analysis Procedure
2. **Randomized and Designed Experiments** allowing to both:
  - Check the model and Instantiate it
3. Careful **recording of all experimental parameters** (before and during XPs)

# REPRODUCIBLE RESEARCH = RIGOR AND TRANSPARENCY

To err is human.

Good research requires time and resources

1. **Train yourself and your students:** RR, statistics, experiments
  - Beware of checklists and norms
  - Understand what's at stake
2. **Change the norm:** make publication practices evolve
3. **Incentive:** consider RR/open science when hiring/promoting

# REPRODUCIBLE RESEARCH = RIGOR AND TRANSPARENCY

To err is human.

Good research requires time and resources

1. **Train yourself and your students:** RR, statistics, experiments
  - Beware of checklists and norms
  - Understand what's at stake
2. **Change the norm:** make publication practices evolve
3. **Incentive:** consider RR/open science when hiring/promoting
4. **Prepare the Future:** Toward **literate experimentation?**
  - Reuse, reuse, reuse!
  - Shared and controlled testbeds (e.g., Grid'5000/SILECS)
  - How to share Experiments ?



# SOME ADVERTISING

**TEN YEARS REPRODUCIBILITY CHALLENGE**

RESCIENCE SPECIAL ISSUE  
FREE TO READ - FREE TO PUBLISH



Would you dare to run the code from your past self?  
(the one that does not answer mail)

SUBMISSION DEADLINE 01/04/2020  
<http://rescience.github.io/ten-years>  
In association with Inria, CNRS, Software Heritage, Rescience, Comité pour la Science Ouverte, DFRST Bordeaux & Mission de la pédagogie et du numérique pour l'enseignement supérieur.



3rd Edition: ≈ Feb. 2020

A new MOOC: "Advanced R"

- Software environment control (Docker)
- Scientific workflow (snakemake)
- Managing data (HDF5, archiving)

October 2020 ?

<http://rescience.github.io/ten-years/>

11/100